


■ Default + Belady
 ■ Default + ILP-Greedy
 ■ COSMA (native ILP)
 ■ Baseline FAILS outright (allocator fragmentation)


degenerate $M_R = \text{MPMF}$ -- open, see text

ResNet-50 (FP32)




degenerate $M_R = \text{MPMF}$ -- open, see text

ResNet-50 (INT8)




degenerate + 513² input -- open, see text

DeepLabV3 (FP32)




degenerate $M_R = \text{MPMF}$ -- open, see text

S3D (FP32)




matches paper's named exception

FCN (FP32)




matches paper's pattern

DenseNet-121 (INT8)




matches paper's pattern

ResNeXt-50 (FP32)



matches at M_P only, see text

R2Plus1D-18 (FP32)



Top row (ResNet-50 x2, DeepLabV3, S3D): structurally degenerate in every toolchain we tested (3 independent exporters) -- $M_R = \text{MPMF}$ with zero budget slack, so every scheme ties at 0. Contradicts the paper's own Fig. 3, which shows nonzero baselines for these same models even at M_P -- investigated at length, not yet resolved.
 Bottom row: FCN reproduces the paper's one named exception; DenseNet-121/ResNeXt-50/R2Plus1D-18 reproduce the paper's actual headline comparison (baseline fails or lags, COSMA wins) on a real, non-degenerate budget.